

C9W5 · THURSDAY, FEBRUARY 19, 2026

BRAND POSITION + SYNTHETIC USER TESTING

From Onboarding Docs to Copy That Converts

PART 1: SYNTHESIS REFLECTION

**WHAT DID YOU LEARN ABOUT AI
BY DOING THE SYNTHESIS YOURSELF
FIRST?**

| DISCUSSION

- What did the AI find that **you missed**?
- What did you notice that the **AI didn't**?
- Where did you **trust** the AI's read? Where didn't you?



2-3 teams share highlights. 2-3 min each.

| THE BRIDGE TO SYNTHETIC TESTING

LAST CLASS

Humans & AI synthesize human interview data

n=10

THIS CLASS

Human output becomes an input for AI to simulate more human responses. AI generates reports for human review.

n=60+

Same critical question: where do you trust it?

PART 2: AI CONTEXT, BRAND POSITION + SYNTHETIC TESTING

AI DEVELOPMENT IS CONTEXT-DRIVEN.

Every new chat window starts from zero.

| THE CONTEXT PROBLEM

- › Every new context = a new "employee" who knows **nothing** about your project
- › Long contexts **degrade**: the AI contradicts itself, forgets decisions, gets fixated
- › So the question becomes: **how do you onboard a new AI quickly?**



Imagine hiring a new contractor **every single session**, one who's never seen your codebase or brand guidelines. You'd write really good onboarding docs.

| AGENT ONBOARDING

BRIEF.MD

What you're building, who it's for, the core problem
Create after interviews + synthesis. Update after major changes.

BRAND_POSITION.MD

How your brand talks: thesis, canonical language, tone, guardrails
Create after interviews + synthesis. Review after any real user feedback.

ARCHITECTURE.MD

How your codebase is set up, major modules, where things live
Create when you start app development. Review after major feature development.

STYLE_GUIDE.MD

Colors, typography, component patterns, visual identity
Create before homepage build. Update when you make components you want to reuse



Today we're building the first two. Architecture and style guide come in Weeks 7-8.

| WHY THIS MATTERS RIGHT NOW



SYNTHETIC TESTING WITH NO CONTEXT DOCS

Generic AI personas giving generic feedback on generic copy



SYNTHETIC TESTING WITH BRIEF + BRAND_POSITION

Grounded personas giving useful feedback on coherent copy

| BRAND POSITION: COPY COMPASS

An **internal** document your whole team agrees on:

- › Defines your thesis, canonical language, tone, say-don't-say
- › All customer-facing copy flows from this

TRADITIONAL

- › Aligns a team of copywriters
- › New hire reads it, internalizes over weeks
- › Updated quarterly
- › Nice to have

AI DEVELOPMENT

- › Aligns **every AI context** you open
- › New AI reads it instantly, applies immediately
- › Referenced **as-needed** based on LLM discretion (and CLAUDE.MD guidance)
- › **Load-bearing infrastructure**

| KEY SECTIONS

BRAND THESIS

Why your product exists and what changes because it does (1-2 sentences)

CANONICAL LANGUAGE

The exact words and phrases you use

POSITIONING STATEMENT

For [audience] who [need], [product] is the [category] that [key benefit]

TARGET IDENTITY

Who is this for? In **their** language, from interviews

LANGUAGE TO AVOID

Words that undermine your brand



Grab the template from the repo, and **ask the AI to walk you through it like a marketer** and fill it out.

SYNTHETIC USER TESTING

| WHAT IS SYNTHETIC USER TESTING?

- AI-generated personas, built from **your** interview data — not generic, grounded in the real people you talked to
- Each persona reads your landing page and gives structured feedback:
 - – Would they sign up?
 - – What's unclear? What's a dealbreaker?
 - – What language resonates?



This is a **preflight check**, not validation. It generates hypotheses. The landing page with real traffic is the actual test.

| WHY IT WORKS

75%

ALIGNMENT WITH
REAL PANELS

Mars Wrigley + Empathy
Lab

0.85

CORRELATION WITH
REAL EXPERIMENTS
Stanford, 476 treatments

90%

CORRELATION ON
RANKINGS

PyMC Labs, 57 real surveys



Copy testing and concept ranking are the strongest use cases — exactly what you're doing.

I WHY IT FAILS



SYCOPHANCY

Synthetic users are **systematically too nice**.
Politeness markers 39% vs 20% for real
humans.



PERSONA FLATTENING

LLMs **reduce within-group diversity**. Different
personas give suspiciously similar feedback.



CAN'T OBSERVE BEHAVIOR

Can evaluate descriptions but can't **actually use** your product, feel friction, or abandon
checkout.



OUT-GROUP STEREOTYPING

LLMs represent what **outsiders think** of a
group, not what the group thinks of **itself**.

| A NOTE ON STATISTICAL SIGNIFICANCE

With **60 personas** (12 per bucket), we're not powered for formal significance tests.

DETECT A 20PT DIFFERENCE

~90 per segment
e.g. 30% → 50% conversion

DETECT A 15PT DIFFERENCE

~160 per segment
e.g. 30% → 45% conversion

OUR SETUP

12 per segment
Directional signal, not significance

⚠️ A [Columbia stats study](#) (Gelman's lab, Dec 2025) found synthetic data adds only **~15% effective sample size** on top of real data.
"Statistical significance on synthetic data" is a house built on sand.

Takeaway: Use 60 for fast iteration and qualitative insight. Scale to 200-500 if you want quantitative claims. But **real traffic is your real significance test.**

| SAY → DO → PAY

1

SAY: INTERVIEWS

What people *told you* about
their problems

2

~DO: SYNTHETIC TESTING

AI predicts reactions by
audience

3

DO: LANDING PAGE

Will strangers actually *sign*
up?

You've collected "Say" evidence. This weekend: preflight "Do" evidence.
Next Tuesday: "*Do*" evidence.

| HOW IT ACTUALLY WORKS

1

BUILD PERSONAS FROM YOUR INTERVIEWS

Your synthesis segments become persona buckets. Grounded in real quotes, real pain points.

2

WRITE V1 COPY FROM YOUR BRAND POSITION

Headlines from interview language. Value props from evidence. CTA matching commitment level.

3

RUN THE TEST

Each persona evaluates: gut reaction, clarity, dealbreakers, unanswered questions, would they sign up?

4

READ THE RESULTS CRITICALLY

Ignore the praise. Focus on dealbreakers, clarity issues, and where personas disagree.

5

REVISE AND REPEAT

Fix biggest issues → V2
→ test → V3. Skill tracks convergence: when scores plateau, you're done.

| THE ITERATION LOOP

1

BUILD PERSONAS
FROM
INTERVIEWS

2

WRITE COPY
FROM BRAND
POSITION

3

RUN SYNTHETIC
TEST

4

READ FEEDBACK,
REVISE COPY, GO
TO 3

5

DEPLOY BEST
VERSION → REAL
TRAFFIC

V1 → V3

Find the framing. Try big swings,
look for concepts that resonate.

V4 → V6

Hone specific language. Identify
anti-audience segments.

V6+

Further refinement + audience
swaps

PART 3: TEAM WORK

| BRAND POSITION: PRIORITY FIELDS

Fill these in class. Use AI to help draft, and go ahead and feed it your interview analyses as well so it can pick up exact user wording.

BRAND THESIS

Why does your product exist? **1-2 sentences.**

TARGET IDENTITY

Who is this for, **in their words?**

CANONICAL LANGUAGE

5-10 phrases from your interviews that you'll use everywhere

LANGUAGE TO AVOID

Founder-speak, jargon, anything that doesn't sound like your users

POSITIONING STATEMENT

For [audience] who [need],
[product] is the [category] that
[key benefit]

| V1 LANDING PAGE COPY

Draft V1 copy with AI. Common sections:

HERO

Pain point headline + your
reframe
From canonical language

HOW IT WORKS

What is this, concretely?
Features, process, examples

SOCIAL PROOF

Why should I believe you?
Testimonials, numbers, logos

OBJECTION HANDLING

Why people say no, and your
answer
From interview hesitations

CTA

What you're asking them to do
Sign up, waitlist, quiz, etc.



Feed your brand position + interview analyses to AI and ask it to draft copy. Find a landing page you like and use it as a structural reference.

| WHAT IS A SKILL?

A **skill** is a detailed instruction file — a markdown document that tells an AI exactly how to do a complex task.

- Works with **Claude, ChatGPT, Gemini**, or any frontier model
- No special tool or plugin needed — just tell the AI to read the file
- Contains the full workflow: what to do first, what to do next, what the output should look like
- Think of it as a **reusable playbook** that any AI can follow



This is the onboarding problem again. Instead of explaining a complex workflow from scratch every time, you hand the AI a skill file.

| HOW TO RUN SYNTHETIC TESTING

1

DOWNLOAD THE SKILL FOLDER

synthetic-user-testing/ from the class exercise materials (also on course repo)

2

OPEN ANY FRONTIER MODEL

Claude, ChatGPT, Gemini — any will work

3

TELL IT TO READ SKILL.MD

It contains the full workflow: personas, evaluations, iteration

4

FEED IT YOUR PERSONA BUCKETS + V1 COPY

Review results: focus on dealbreakers, clarity, what's unclear

5

REVISE → RUN AGAIN

Use the best version for your landing page

PART 4: WRAP + WEEKEND PLAN

| WHAT'S DUE TUESDAY FEB 24

1. BRAND POSITION

Polished from today's draft

2. V1 LANDING PAGE COPY

Ready for testing

3. SYNTHETIC TESTING

At least 2 rounds of iteration (V1 → V2 minimum, V3 if possible)

4. DEPLOYED LANDING PAGE

Real URL, live on the internet, using your best copy version

Platforms: Carrd, Framer, Webflow, raw HTML — whatever works

Must include: headline, value prop, features/bullets, CTA

| THE FULL ARC



**WEEK 4:
PROBLEM +
SOLUTION
INTERVIEWS**

**WEEK 5 TUE:
INTERVIEW
SYNTHESIS +
DOWNSELECT**

**WEEK 5 THU:
BRAND
POSITION +
SYNTHETIC
TESTING**

Today

**WEEK 6 TUE:
LANDING
PAGES DUE +
TRAFFIC
STRATEGY**

Next class

**WEEK 6 THU:
DEMAND
TEST –
ANALYZE
RESULTS,
ITERATE**

**WEEK 7:
START
BUILDING
THE REAL
PRODUCT**



QUESTIONS?

END OF CLASS 9

Next: Landing Pages Due + Traffic Strategy